

Kubernetes Assignment 01

1. List all the namespaces in the cluster
`kubectl get namespace`
2. List all the pods in all namespaces
`kubectl get pods --all-namespaces`
3. List all the pods in the particular namespace
`kubectl get pods --namespace <nsname>`
4. List all the services in the particular namespace
`kubectl get services --namespace <nsname>`
5. List all the pods showing name and namespace with a json path expression
`kubectl get pods -o=jsonpath="{.items[*]['metadata.name' , 'metadata.namespace']}"`
6. Create an nginx pod in a default namespace and verify the pod running
`kubectl -- run nginx --image=nginx -n default`
7. Create the same nginx pod with a yaml file
`kubectl create -f nginxpod.yaml`
8. Output the yaml file of the pod you just created
`kubectl get deploy nginx-deployment -o yaml`
9. Output the yaml file of the pod you just created without the cluster-specific information
`kubectl get pod nginx -o yaml --export (deprecated)`
10. Get the complete details of the pod you just created
`kubectl describe pod nginx`
11. Delete the pod you just created
`kubectl delete pod nginx`
12. Delete the pod you just created without any delay
`kubectl delete pod nginx --grace-period=0 --force`
13. Create the nginx pod with version 1.17.4 and expose it on port 80
`kubectrl run nginx -- image=nginx-pod -- port=80 - -expose`

14. Change the Image version to 1.15-alpine for the pod you just created and verify the image version is updated
`kubectl set image pod/nginx nginx=nginx:1.15-alpine`
15. Change the Image version back to 1.17.1 for the pod you just updated and observe the change
`kubectl set image pod/nginx nginx=nginx:1.17.1`
`kubectl describe pod nginx`
`kubectl get pod nginx -w`
16. Check the Image version without the describe command
`kubectl get pod nginx -o jsonpath='{.spec.containers[].image}'`
17. Create the nginx pod and execute the simple shell on the pod
`kubectl -- exec -it nginx bash`
18. Get the IP Address of the pod you just created
`kubectl describe pod nginx`
19. Create a busybox pod and run command ls while creating it and check the logs
`kubectl run busybox --image=nginx -n default`
`minikube kubectl logs busybox`
20. If pod crashed check the previous logs of the pod
`kubectl logs --container < container_name > --previous=true < pod_name >`
21. Create a busybox pod with command sleep 3600
`kubectl run busybox --image=busybox --restart=Never -- /bin/sh -c "sleep 3600"`
22. Check the connection of the nginx pod from the busybox pod
`kubectl exec -it busybox -- wget -o- 127.0.0.4`
23. Create a busybox pod and echo message 'How are you' and delete it manually
`kubectl run busybox --image=nginx --restart=Never -it -- echo "How are you"`
`kubectl delete pod busybox`
24. Create a busybox pod and echo message 'How are you' and have it deleted immediately
`kubectl run busybox --image=nginx --restart=Never -it --rm -- echo "How are you"`
25. Create an nginx pod and list the pod with different levels of verbosity
`kubectl run busybox --image=nginx --restart=Never --v=7`
26. List the nginx pod with custom columns POD_NAME and POD_STATUS

```
kubectl get po -o=custom-columns="POD_NAME:.metadata.name,  
POD_STATUS:.status.containerStatuses[].state"
```

27. List all the pods sorted by name

```
kubectl get pods --sort-by=.metadata.name
```

28. List all the pods sorted by created timestamp

```
kubectl get po --sort-by=.status.startTime
```

29. Create a Pod with three busy box containers with commands "ls; sleep 3600;", "echo Hello World; sleep 3600;" and "echo this is the third container; sleep 3600" respectively and check the status

```
Kubectl create -f multic.yaml
```

30. Check the logs of each container that you just created

```
kubectl logs busybox -c busybox1
```

31. Check the previous logs of the second container busybox2 if any

```
kubectl logs busybox -c busybox2 --previous=true
```

32. Run command ls in the third container busybox3 of the above pod

```
kubectl exec busybox -c busybox3 -- ls
```

33. Show metrics of the above pod containers and puts them into the file.log and verify

```
minikube addons enable metrics-server
```

```
kubectl top pod busybox --containers
```

```
kubectl top pod busybox --containers > file.log
```

```
cat file.log
```

34. Create a Pod with main container busybox and which executes this "while true; do echo 'Hi I am from Main container' >> /var/log/index.html; sleep 5; done" and with sidecar container with nginx image which exposes on port 80. Use emptyDir Volume and mount this volume on path /var/log for busybox and on path /usr/share/nginx/html for nginx container. Verify both containers are running.

```
kubectl create -f Q35.yaml
```

```
Kubectl get pods
```

```
kubectl exec -it multi-cont-pod -c main-container -- sh
```

```
cat /var/log/index.html
```

```
kubectl exec -it multi-cont-pod -c sidecar-container -- sh
```

```
cat /usr/share/nginx/html/index.html
```

35. Exec into both containers and verify that main.txt exist and query the main.txt from sidecar container with curl localhost

```
kubectl exec -it multi-cont-pod -c sidecar-container -- /bin/bash
```

`curl localhost`

36. Get the pods with label information

`kubectl get pods --show-labels`

37. Create 5 nginx pods in which two of them is labeled env=prod and three of them is labeled env=dev

`kubectl run nginx-prod1 --image=nginx --restart=Never --labels=env=prod`

`kubectl run nginx-prod2 --image=nginx --restart=Never --labels=env=prod`

`kubectl run nginx-dev1 --image=nginx --restart=Never --labels=env=dev`

`kubectl run nginx-dev2 --image=nginx --restart=Never --labels=env=dev`

`kubectl run nginx-dev3 --image=nginx --restart=Never --labels=env=dev`

38. Verify all the pods are created with correct labels

`kubectl get pods --show-labels`

39. Get the pods with label env=dev

`kubectl get pods -l env=dev`

40. Get the pods with label env=dev and also output the labels

`kubectl get pods -l env=dev --show-labels`

41. Get the pods with label env=prod

`kubectl get pods -l env=prod`

42. Get the pods with label env=prod and also output the labels

`kubectl get pods -l env=prod --show-labels`

43. Get the pods with label env

`kubectl get pods -l env`

44. Get the pods with labels env=dev and env=prod

`kubectl get pods -l 'env in (dev,prod)'`

45. Get the pods with labels env=dev and env=prod and output the labels as well

`kubectl get pods -l 'env in (dev,prod)' --show-labels`

46. Change the label for one of the pod to env=uat and list all the pods to verify

`kubectl label pod/nginx-dev1 env=uat --overwrite`

47. Remove the labels for the pods that we created now and verify all the labels are removed

`kubectl label pod nginx-dev{1..3} env-`

`kubectl label pod nginx-prod{1..2} env-`

```
kubectl get pods --show-labels
```

48. Let's add the label app=nginx for all the pods and verify

```
kubectl label pod nginx-dev{1..3} app=nginx  
kubectl label pod nginx-prod{1..2} app=nginx  
kubectl get pod --show-labels
```

49. Get all the nodes with labels (if using minikube you would get only master node)

```
kubectl get nodes --show-labels
```

50. Label the node (minikube if you are using) nodeName=nginxnode

```
kubectl label node minikube nodeName=nginxnode
```

51. Create a Pod that will be deployed on this node with the label nodeName=nginxnode

[Bbyaml.yaml](#)

52. Verify the pod that it is scheduled with the node selector

```
kubectl describe pod nginx | grep Node-Selectors
```

53. Verify the pod nginx that we just created has this label

```
kubectl describe pod nginx
```

54. Annotate the pods with name=webapp

```
kubectl annotate pod nginx-dev{1..3} name=webapp  
kubectl annotate pod nginx-prod{1..2} name=webapp
```

55. Verify the pods that have been annotated correctly

```
kubectl describe pod nginx-dev{1..3} | grep -i annotations  
kubectl describe pod nginx-prod{1..2} | grep -i annotations
```

56. Remove the annotations on the pods and verify

```
kubectl annotate pod nginx-dev{1..3} name-  
kubectl annotate pod nginx-prod{1..2} name-  
kubectl describe po nginx-dev{1..3} | grep -i annotations  
kubectl describe po nginx-prod{1..2} | grep -i annotations
```

57. Remove all the pods that we created so far

```
kubectl delete pod --all
```

58. Create a deployment called webapp with image nginx with 5 replicas

```
kubectl create -f Q58.yaml
```

59. Get the deployment you just created with labels

```
kubectl get deploy webapp --show-labels
```

60. Output the yaml file of the deployment you just created
`kubectl get deploy webapp -o yaml`
61. Get the pods of this deployment
`kubectl get pod -l app=webapp`
62. Scale the deployment from 5 replicas to 20 replicas and verify
`kubectl scale deploy webapp --replicas=20`
`kubectl get rs -l app=webapp`
63. Get the deployment rollout status
`kubectl rollout status deploy webapp`
64. Get the replicaset that created with this deployment
`kubectl get rs -l app=webapp`
65. Get the yaml of the replicaset and pods of this deployment
`kubectl get rs -l app=webapp -o yaml`
`kubectl get pod -l app=webapp -o yaml`
66. Delete the deployment you just created and watch all the pods are also being deleted
`kubectl delete deploy webapp`
`kubectl get pod -l app=webapp -w`
67. Create a deployment of webapp with image nginx:1.17.1 with container port 80 and verify the image version
`kubectl create -f wapp.yaml`
`kubectl describe deploy webapp | grep Image`
68. Update the deployment with the image version 1.17.4 and verify
`kubectl set image deploy/webapp nginx=nginx:1.17.4`
`kubectl describe deploy webapp | grep Image`
69. Check the rollout history and make sure everything is ok after the update
`kubectl rollout history deploy webapp`
70. Undo the deployment to the previous version 1.17.1 and verify Image has the previous version
`kubectl rollout undo deploy/webapp`
`kubectl describe deploy webapp | grep Image`
71. Update the deployment with the image version 1.16.1 and verify the image and also check the rollout history
`kubectl set image deploy/webapp nginx=nginx:1.16.1`

```
kubectl describe deploy webapp | grep Image  
kubectl rollout history deploy webapp
```

72. Update the deployment to the Image 1.17.1 and verify everything is ok

```
kubectl rollout undo deploy webapp --to-revision=3
```

73. Update the deployment with the wrong image version 1.100 and verify something is wrong with the deployment

```
kubectl set image deploy/webapp nginx=nginx:1.100 (imagepullbackoff status)  
kubectl rollout status deploy webapp
```

74. Undo the deployment with the previous version and verify everything is Ok

```
kubectl rollout undo deploy webapp  
kubectl rollout status deploy webapp  
kubectl get pod -l app=webapp
```

75. Check the history of the specific revision of that deployment

```
kubectl rollout history deploy webapp --revision=7
```

76. Pause the rollout of the deployment

```
kubectl rollout pause deploy webapp
```

77. Update the deployment with the image version latest and check the history and verify nothing is going on

```
kubectl set image deploy webapp nginx=nginx:latest  
kubectl rollout history deploy webapp
```

78. Resume the rollout of the deployment

```
kubectl rollout resume deploy webapp
```

79. Check the rollout history and verify it has the new version

```
kubectl rollout history deploy webapp
```

80. Apply the autoscaling to this deployment with minimum 10 and maximum 20 replicas and target CPU of 85% and verify hpa is created and replicas are increased to 10 from 1

```
kubectl autoscale deploy webapp --min=1 --max=10 --cpu-percent=85  
kubectl get hpa
```

81. Clean the cluster by deleting deployment and hpa you just created

```
kubectl delete deploy webapp  
kubectl delete hpa webapp
```

82. Create a Job with an image node which prints node version and also verifies there is a pod created for this job

```
kubectrl create job myjob --image=node -- node -v
kubectrl get pod
```

83. Get the logs of the job just created

```
kubectrl logs myjob-knmr2
```

84. Output the yaml file for the Job with the image busybox which echos "Hello I am from job"

```
kubectrl create job hello-job --image=busybox --dry-run -o yaml -- echo "Hello I am from job"
```

85. Copy the above YAML file to hello-job.yaml file and create the job

```
kubectrl create job job2 --image=busybox --dry-run=client -o yaml -- echo "Hey there" >
hello-job.yaml
kubectrl create -f hello-job.yaml
```

86. Verify the job and the associated pod is created and check the logs as well

```
kubectrl get job
kubectrl get po
kubectrl logs job2-kj8q4
```

87. Delete the job we just created

```
kubectrl delete job job2
```

88. Create the same job and make it run 10 times one after one

```
kubectrl create -f hello-job.yaml
```

89. Watch the job that runs 10 times one by one and verify 10 pods are created and delete those after it's completed

```
kubectrl get job -w
Kubectrl get pods
kubectrl delete job job2
```

90. Create the same job and make it run 10 times parallel

Add parallelism: 10 in spec of yaml

91. Watch the job that runs 10 times parallelly and verify 10 pods are created and delete those after it's completed

```
kubectrl get job -w
kubectrl delete job job2
```

92. Create a Cronjob with busybox image that prints date and hello from kubernetes cluster message for every minute


```
kubectrl create cronjob democronjob --image=busybox --schedule="*/1 * * * *" -- bin/sh -c  
"date; echo Hello from kubernetes cluster"
```

93. Output the YAML file of the above cronjob

```
kubectrl get cj democronjob -o yaml
```

94. Verify that CronJob creating a separate job and pods for every minute to run and verify the logs of the pod

```
Kubectrl get pod
```

```
kubectrl logs democronjob-1613996340-lhn54
```

95. Delete the CronJob and verify all the associated jobs and pods are also deleted.

```
kubectrl delete cj democronjob
```

```
Kubectrl get pod
```

```
Kubectrl get job
```

96. List Persistent Volumes in the cluster

```
kubectrl get pv
```

97. Create a hostPath PersistentVolume named task-pv-volume with storage 10Gi, access modes ReadWriteOnce, storageClassName manual, and volume at /mnt/data and verify

```
kubectrl create -f task-pv-volume.yaml
```

```
kubectrl get pv
```

98. Create a PersistentVolumeClaim of at least 3Gi storage and access mode ReadWriteOnce and verify status is Bound

```
kubectrl create -f task-pv-claim.yaml
```

```
kubectrl get pvc
```

99. Delete persistent volume and PersistentVolumeClaim we just created

```
Kubectrl delete pv --all
```

```
Kubectrl delete pvc --all
```

100. Create a Pod with an image Redis and configure a volume that lasts for the lifetime of the Pod

```
kubectrl create -f Q100.yaml
```

101. Exec into the above pod and create a file named file.txt with the text 'This is called the file' in the path /data/redis and open another tab and exec again with the same pod and verifies file exist in the same path.

```
kubectrl exec -it redis-pod -- bash
```

```
cd /data/redis
```

```
Touch file.txt
```

```
echo 'This is called the file' > file.txt
```

```
New tab --> kubectl exec -it redis-pod -- bash
cat /data/redis/file.txt
```

102. Delete the above pod and create again from the same yaml file and verifies there is no file.txt in the path /data/redis

```
kubectl delete pod redis-pod
kubectl create -f Q100.yaml
kubectl exec -it redis-pod -- bash
cat /data/redis/file.txt
```

103. Create PersistentVolume named task-pv-volume with storage 10Gi, access modes ReadWriteOnce, storageClassName manual, and volume at /mnt/data and Create a PersistentVolumeClaim of at least 3Gi storage and access mode ReadWriteOnce and verify status is Bound

```
kubectl create -f demopv.yaml
kubectl create -f Q98.yaml
```

104. Create an nginx pod with containerPort 80 and with a PersistentVolumeClaim task-pv-claim and has a mount path "/usr/share/nginx/html"

```
kubectl create -f Q104.yaml
```

105. List all the configmaps in the cluster

```
kubectl get cm
```

106. Create a configmap called myconfigmap with literal value appname=myapp

```
kubectl create cm myconfigmap --from-literal=appname=myapp
```

107. Verify the configmap we just created has this data

```
kubectl get cm -o yaml
```

108. delete the configmap myconfigmap we just created

```
Kubectl delete cm myconfigmap
```

109. Create a file called config.txt with two values key1=value1 and key2=value2 and verify the file

```
cat >> config.txt << EOF
key1=value1
key2=value2
EOF
cat config.txt
```

110. Create a configmap named keyvalcfgmap and read data from the file config.txt and verify that configmap is created correctly

```
kubectl create cm keyvalcfgmap --from-file=config.txt
```

```
kubectl get cm keyvalcfgmap -o yaml
```

111. Create an nginx pod and load environment values from the above configmap keyvalcfgmap and exec into the pod and verify the environment variables and delete the pod

```
kubectl create -f Q111.yaml  
kubectl exec -it nginx-pod -- env  
kubectl delete pod nginx-pod
```

112. Create an env file file.env with var1=val1 and create a configmap envcfgmap from this env file and verify the configmap

```
kubectl create cm envcfgmap --from-env-file=file.env  
kubectl get cm envcfgmap -o yaml
```

113. Create an nginx pod and load environment values from the above configmap envcfgmap and exec into the pod and verify the environment variables and delete the pod

```
kubectl create -f Q113.yaml  
kubectl exec -it nginx -- env  
kubectl delete pod nginx
```

114. Create a configmap called cfgvolume with values var1=val1, var2=val2 and create an nginx pod with volume nginx-volume which reads data from this configmap cfgvolume and put it on the path /etc/cfg

```
kubectl create cm cfgvolume --from-literal=var1=val1 --from-literal=var2=val2  
kubectl describe cm cfgvolume  
kubectl create -f Q114.yaml  
Terminal: kubectl exec -it nginxpod -- bin/sh  
# cd /etc/cfg  
# ls  
var1 var2
```

115. Create a pod called secbusybox with the image busybox which executes command sleep 3600 and makes sure any Containers in the Pod, all processes run with user ID 1000 and with group id 2000 and verify.

```
kubectl create -f Q115.yaml  
kubectl run secbusybox --image=busybox --restart=Never --dry-run -o yaml -- /bin/sh -c  
"sleep 3600;" > Q115.yaml  
kubectl exec -it secbusybox -- sh  
$id  
Output : uid=1000 gid=2000 groups=2000
```

116. Create the same pod as above this time set the securityContext for the container as well and verify that the securityContext of container overrides the Pod level securityContext.

```
kubectl create -f Q116.yaml
kubectl exec -it secbusybox1 -- sh
$id → uid:2000
```

117. Create pod with an nginx image and configure the pod with capabilities NET_ADMIN and SYS_TIME verify the capabilities
kubectl create -f Q117.yaml

```
[ayrastog@ayrastog nginxpod]$ kubectl exec -it nginxpod -- sh
# cd /proc/1
# cat status
Name:      nginx
Umask:    0022
State:     S (sleeping)
Tgid:     1
Ngid:     0
Pid:      1
PPid:     0
TracerPid: 0
Uid:      0      0      0      0
Gid:      0      0      0      0
FDSize:   64
Groups:
NSTgid:   1
NSpid:    1
NSpgid:   1
NSSid:    1
VmPeak:   10664 kB
VmSize:   10636 kB
VmLck:     0 kB
VmPin:     0 kB
VmHWM:     5876 kB
VmRSS:     5876 kB
RssAnon:           804 kB
RssFile:          5072 kB
RssShmem:           0 kB
VmData:     980 kB
VmStk:      132 kB
VmExe:      980 kB
VmLib:     3800 kB
VmPTE:       60 kB
VmSwap:       0 kB
HugetlbPages:      0 kB
CoreDumping: 0
Threads:      1
SigQ:  0/22437
SigPnd: 0000000000000000
ShdPnd: 0000000000000000
SigBlk: 0000000000000000
SigIgn: 0000000040001000
SigCgt: 0000000198016a07
CapInh: 00000000aa0435fb
CapPrm: 00000000aa0435fb
CapEff: 00000000aa0435fb
CapBnd: 00000000aa0435fb
CapAmb: 0000000000000000
NoNewPrivs: 0
Seccomp: 0
Speculation_Store_Bypass: thread vulnerable
Cpus_allowed: 3
Cpus_allowed_list: 0-1
```

118. Create a Pod nginx and specify a memory request and a memory limit of 100Mi and 200Mi respectively.

`kubectrl create -f Q118.yaml`

119. Create a Pod nginx and specify a CPU request and a CPU limit of 0.5 and 1 respectively.

`kubectrl create -f Q119.yaml`

120. Create a Pod nginx and specify both CPU, memory requests and limits together and verify.

`kubectrl create -f Q120.yaml`

`kubectrl top pod`

121. Create a Pod nginx and specify a memory request and a memory limit of 100Gi and 200Gi respectively which is too big for the nodes and verify pod fails to start because of insufficient memory

`kubectrl create -f Q121.yml`

`kubectrl describe pod nginx - pending state`

122. Create a secret mysecret with values user=myuser and password=mypassword

`kubectrl create secret generic my-secret --from-literal=user=myuser`

`--from-literal=password=mypassword`

123. List the secrets in all namespaces

`kubectrl get secret --all-namespaces`

124. Output the yaml of the secret created above

`kubectrl get secret my-secret -o yaml`

125. Create an nginx pod which reads username as the environment variable

`kubectrl create -f Q125.yaml`

126. Create an nginx pod which loads the secret as environment variables

`kubectrl create -f Q126.yaml`

`Kubectrl exec -it nginx -- env`

127. List all the service accounts in the default namespace

`kubectrl get sa`

128. List all the service accounts in all namespaces

`kubectrl get sa --all-namespaces`

129. Create a service account called admin

`kubectrl create sa admin`

130. Output the YAML file for the service account we just created

`Kubectl get sa admin -o yaml`

131. Create a busybox pod which executes this command sleep 3600 with the service account admin and verify

`kubectl create -f Q131.yaml`

`kubectl describe pod busyboxpod`

132. Create an nginx pod with containerPort 80 and it should only receive traffic only it checks the endpoint / on port 80 and verify and delete the pod.

`kubectl create -f nginx-get.yaml`

`kubectl describe pod nginx | grep -i readiness`

`kubectl delete pod nginx`

133. Create an nginx pod with containerPort 80 and it should check the pod running at endpoint / healthz on port 80 and verify and delete the pod.

`Change path to /healthz`

134. Create an nginx pod with containerPort 80 and it should check the pod running at endpoint /healthz on port 80 and it should only receive traffic only it checks the endpoint / on port 80. verify the pod.

`Add readinessProbe`

135. Check what all are the options that we can configure with readiness and liveness probes

`kubectl explain Pod.spec.containers.livenessProbe`

`kubectl explain Pod.spec.containers.readinessProbe`

136. Create the pod nginx with the above liveness and readiness probes so that it should wait for 20 seconds before it checks liveness and readiness probes and it should check every 25 seconds.

`kubectl create -f Q136.yaml`

137. Create a busybox pod with this command "echo I am from busybox pod; sleep 3600;" and verify the logs.

`kubectl run busybox --image=busybox --restart=Never -- /bin/sh -c "echo I am from busybox pod; sleep 3600;"`

`kubectl logs busybox`

138. copy the logs of the above pod to the busybox-logs.txt and verify

`kubectl logs busybox > busybox-logs.txt`

`cat busybox-logs.txt`

139. List all the events sorted by timestamp and put them into file.log and verify

`kubectl get events --sort-by=.metadata.creationTimestamp`

```
kubectl get events --sort-by=.metadata.creationTimestamp > file.log
cat file.log
```

140. Create a pod with an image alpine which executes this command "while true; do echo 'Hi I am from alpine'; sleep 5; done" and verify and follow the logs of the pod.

```
kubectl run hello --image=alpine --restart=Never -- /bin/sh -c "while true; do echo 'Hi I
am from Alpine'; sleep 5;done"
kubectl logs --follow hello
```

141. Create the pod with this kubectl create -f

```
https://gist.githubusercontent.com/bbachi/212168375b39e36e2e2984c097167b00/raw/1f
d63509c3ae3a3d3da844640fb4cca744543c1c/not-running.yml. The pod is not in the
running state. Debug it.
```

```
kubectl create -f
https://gist.githubusercontent.com/bbachi/212168375b39e36e2e2984c097167b00/raw/1f
d63509c3ae3a3d3da844640fb4cca744543c1c/not-running.yml
kubectl get pod not-running
kubectl describe pod not-running - ImagePullBackOff
kubectl set image pod/not-running not-running=nginx
```

142. This following yaml creates 4 namespaces and 4 pods. One of the pod in one of the namespaces are not in the running state. Debug and fix it.

```
https://gist.githubusercontent.com/bbachi/1f001f10337234d46806929d12245397/raw/84
b7295fb077f15de979fec5b3f7a13fc69c6d83/problem-pod.yaml.
```

```
kubectl create -f
https://gist.githubusercontent.com/bbachi/1f001f10337234d46806929d12245397/raw/84
b7295fb077f15de979fec5b3f7a13fc69c6d83/problem-pod.yaml
kubectl get po --all-namespaces
kubectl get po -n namespace2
kubectl set image pod/pod2 pod2=nginx -n namespace2
kubectl get po -n namespace2
```

143. Get the memory and CPU usage of all the pods and find out top 3 pods which have the highest usage and put them into the cpu-usage.txt file

```
kubectl top pod --all-namespaces | sort --reverse --key 3 --numeric | head -3 >
cpu-usage.txt
cat cpu-usage.txt
```

144. Create an nginx pod with a yaml file with label my-nginx and expose the port 80

```
kubectl run nginx --image=nginx --restart=Never --port=80 --dry-run -o yaml > Q144.yaml
kubectl create -f Q144.yaml
```

145. Create the service for this nginx pod with the pod selector app: my-nginx

```
kubectl create -f Q145.yaml
```

146. Find out the label of the pod and verify the service has the same label
- ```
kubectl get pod nginx --show-labels
kubectl get svc my-service -o wide
```
147. Delete the service and create the service with kubectl expose command and verify the label
- ```
kubectl delete svc my-service
kubectl expose po nginx --port=80 --target-port=9376
kubectl get svc -l app=my-nginx
```
148. Delete the service and create the service again with type NodePort
- ```
kubectl delete svc nginx
kubectl expose po nginx --port=80 --type=NodePort
```
149. Create the temporary busybox pod and hit the service. Verify the service that it should return the nginx page index.html.
- ```
kubectl get svc nginx -o wide
kubectl run busybox --image=busybox --restart=Never -it --rm -- wget -o- <Cluster IP>:80
```
150. Create a NetworkPolicy which denies all ingress traffic
- ```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
Metadata:
 name: default-deny
spec:
 podSelector: {}
 policyTypes:
 - Ingress
```